

# Gradient Descent Method in Solving Convex Optimization Problems

Presenter: Siyuan Lin & Gregory Matesi  
Department of Mathematical and Statistical Sciences  
University of Colorado, Denver  
Fall 2019

# General Introduction


- **Taylor Series:** Let  $f(x)$  be any function infinitely differentiable around  $x = x_0$ .

$$\begin{aligned} f(x) &= \sum_{k=0}^{\infty} \frac{f^{(n)}(x_0)}{n!} (x - x_0)^n \\ &= f(x_0) + f'(x_0)(x - x_0) + \frac{f''(x_0)}{2!} (x - x_0)^2 + \dots \end{aligned}$$

- When  $x$  is close to  $x_0$    $f(x) \approx f(x_0) + f'(x_0)(x - x_0)$

E.g. Taylor series for  $h(x)=\sin(x)$  around  $x_0=\pi/4$

$$\sin(x) = \frac{1}{\sqrt{2}} + \frac{x - \frac{\pi}{4}}{\sqrt{2}} - \frac{\left(x - \frac{\pi}{4}\right)^2}{2\sqrt{2}} - \frac{\left(x - \frac{\pi}{4}\right)^3}{6\sqrt{2}} + \frac{\left(x - \frac{\pi}{4}\right)^4}{24\sqrt{2}} + \frac{\left(x - \frac{\pi}{4}\right)^5}{120\sqrt{2}} - \frac{\left(x - \frac{\pi}{4}\right)^6}{720\sqrt{2}} - \frac{\left(x - \frac{\pi}{4}\right)^7}{5040\sqrt{2}} + \frac{\left(x - \frac{\pi}{4}\right)^8}{40320\sqrt{2}} + \frac{\left(x - \frac{\pi}{4}\right)^9}{362880\sqrt{2}} - \frac{\left(x - \frac{\pi}{4}\right)^{10}}{3628800\sqrt{2}} + \dots$$



The approximation is good around  $\frac{\pi}{4}$ .

# Multivariable Taylor Series

$$\blacktriangleright f(x, y) = f(x_0, y_0) + \frac{\partial f(x_0, y_0)}{\partial x} (x - x_0) + \frac{\partial f(x_0, y_0)}{\partial y} (y - y_0) \\ + \text{something related to } (x - x_0)^2 \text{ and } (y - y_0)^2 + \dots$$

When  $x$  and  $y$  is close to  $x_0$  and  $y_0$



$$f(x, y) \approx f(x_0, y_0) + \frac{\partial f(x_0, y_0)}{\partial x} (x - x_0) + \frac{\partial f(x_0, y_0)}{\partial y} (y - y_0)$$

Consider the second order term, e.g. Newton's Method

# Formal Derivation

## ► Target Function

$$\min_x f(x)$$

## ► Gradient

$$\nabla f(x_0) = f'(x_0)$$

$-\nabla f(x_n)$  : max-rate descending direction


## ► Gradient Descent algorithm

$$x_{n+1} = x_n - \alpha_n \cdot \nabla f(x_n)$$

## ► Step Size

$\alpha_n$  (depends on the assumption of  $f$ )

Learning rate in machine learning




# Pros of gradient descent

- ▶ Simple and intuitive;
  - Work under very few assumptions;
  - Work together with many other methods;
- ▶ Can be very fast for smooth objective functions,
  - i.e. well-conditioned and strongly convex.
  - when the function is convex:
    - ▶ local minima is global minima
    - ▶ converge to the global solution.

# Limitations of Gradient Descent Method

## Zigzagging Issue


- ▶ When target function is non-convex
  - gradient descent increasingly 'zigzags'
  - Newton's method is a better approach.



# Limitations of Gradient Descent Method

## ► Non-convergence Issue

- too small  $\rightarrow$  slow search
- too big  $\rightarrow$  cause overshooting and diverged iterations






# Limitations - Linear Objective function




- Project gradient into affine subspace
  - $A\Delta x = 0$
- Project gradient into first orthant
  - $c^T \Delta x \geq 0$



# Limitations - Linear Objective function



- ▶ Gradient descent also does not perform well with a linear objective function
  - Projection of the gradient onto affine subspace
  - Projecting the gradient into the 1<sup>st</sup> orthant
    - ▶ Computationally expensive
  - Possible infeasible improving direction.



# Newton Methods



## ► Rootfinder


- $f'(x) = 0$

## ► Hessian

- Square matrix of second order partial derivatives
- Curvature information

## ► Quadratic convergence

- In a neighborhood around a local optimum



# Newton Methods



- Derivative of second order Taylor expansion

$$f_T(x) = f_T(x_n + \Delta x) \approx f(x_n) + f'(x_n)\Delta x + \frac{1}{2}f''(x_n)\Delta x^2$$

$$0 = \frac{d}{d\Delta x} \left( f(x_n) + f'(x_n)\Delta x + \frac{1}{2}f''(x_n)\Delta x^2 \right) = f'(x_n) + f''(x_n)\Delta x$$

$$0 = f'(x_n) + f''(x_n)\Delta x$$

# Newton Methods



- Solve for  $\Delta x$

$$\Delta x = -\frac{f'(x_n)}{f''(x_n)}$$

- Iterative process

$$x_{n+1} = x_n + \Delta x = x_n - \frac{f'(x_n)}{f''(x_n)}$$

- Multivariate

$$x_{n+1} = x_n - \gamma [Hf(x_n)]^{-1} \nabla f(x_n)$$
$$\gamma \in (0,1), \gamma = 1$$

# Newton Methods

---



- ▶ Inverting **Hessian** can be computationally expensive
- ▶ Sometimes a surrogate will be used
  - **Semi-newton methods**

# Implementation

---



- ▶ Most solvers use **semi-newton methods**
- ▶ **CONOPT** solver in AMPL.
- ▶ Generalized Reduced Gradient (**GRG**) algorithm
  - **Newton method** option
  - Sequential Quadratic Programming (**SQP**)
  - <https://ampl.com/SOLVERS/conopt3.pdf>

# CONOPT

---



- Takes advantage of appearance of **local linearity**
  - when the function and constraint values are close to their linear approximations
  - Does not compute **Jacobian** for every iteration





## ► Lagrangian

- Linear combination of the objective function with the constraints

$$L(x, \lambda) = f(x) + \lambda g(x)$$

- Hessian of the lagrangian multiplied by a vector

# CONOPT

---



- ▶ nf: number of objective function evaluations
- ▶ ng: number of gradient evaluations
- ▶ nc: number of constraint evaluations
  
- ▶ nJ: number of evaluations of the Jacobian of the constraints
- ▶ nH: number of evaluations of the Hessian of the Lagrangian function
- ▶ **nHv: number of evaluations of the Hessian of the Lagrangian multiplied by a vector**

# Example



► `option solver conopt, conopt_options 'lsesqp=1';`

Logical Switch Enabling SQP mode

`conopt_options 'outlev=2';`

CONOPT 3.17A: outlev=2

The model has 5 variables and 2 constraints  
with 10 Jacobian elements, 5 of which are nonlinear.  
The Hessian of the Lagrangian has 5 elements on the diagonal,  
0 elements below the diagonal, and 5 nonlinear variables.

Pre-triangular equations: 0  
Post-triangular equations: 1

# Example



```
ampl: reset; model mixtures.mod; data mixtures.dat; option solver conopt, conopt_options 'lsesqp=0'; solve; display pi;
CONOPT 3.17A: lsesqp=0
CONOPT 3.17A: Locally optimal; objective 19.64638424
11 iterations; evals: nf = 18, ng = 11, nc = 18, nJ = 0, nH = 0, nHv = 0
pi [*] :=
    african    0.104805
    east_asian 0.264861
    european   0.17273
native_american 0.284324
    south_asian 0.173279
;

ampl:
ampl: reset; model mixtures.mod; data mixtures.dat; option solver conopt, conopt_options 'lsesqp=1'; solve; display pi;
CONOPT 3.17A: lsesqp=1
CONOPT 3.17A: Locally optimal; objective 19.64638424
10 iterations; evals: nf = 15, ng = 9, nc = 15, nJ = 0, nH = 0, nHv = 3
pi [*] :=
    african    0.104805
    east_asian 0.264861
    european   0.17273
native_american 0.284324
    south_asian 0.173279
;
```

# Example



```
ampl: reset; model mixtures.mod; data mixtures.dat; option solver conopt, conopt_options 'lsesqp=0'; solve; display pi;
CONOPT 3.17A: lsesqp=0
CONOPT 3.17A: Locally optimal; objective 19.64638424
11 iterations; evals: nf = 18, ng = 11, nc = 18, nJ = 0, nH = 0, nHv = 0
pi [*] :=
    african    0.104805
    east_asian 0.264861
    european   0.17273
native_american 0.284324
    south_asian 0.173279
;

ampl:
ampl: reset; model mixtures.mod; data mixtures.dat; option solver conopt, conopt_options 'lsesqp=1'; solve; display pi;
CONOPT 3.17A: lsesqp=1
CONOPT 3.17A: Locally optimal; objective 19.64638424
10 iterations; evals: nf = 15, ng = 9, nc = 15, nJ = 0, nH = 0, nHv = 3
pi [*] :=
    african    0.104805
    east_asian 0.264861
    european   0.17273
native_american 0.284324
    south_asian 0.173279
;
```

# Example



```
ampl: reset; model mixtures.mod; data mixtures.dat; option solver conopt, conopt_options 'lsesqp=0'; solve; display pi;
CONOPT 3.17A: lsesqp=0
CONOPT 3.17A: Locally optimal; objective 19.64638424
11 iterations; evals: nf = 18, ng = 11, nc = 18, nJ = 0, nH = 0, nHv = 0
pi [*] :=
    african    0.104805
    east_asian 0.264861
    european   0.17273
native_american 0.284324
    south_asian 0.173279
;

ampl:
ampl: reset; model mixtures.mod; data mixtures.dat; option solver conopt, conopt_options 'lsesqp=1'; solve; display pi;
CONOPT 3.17A: lsesqp=1
CONOPT 3.17A: Locally optimal; objective 19.64638424
10 iterations; evals: nf = 15, ng = 9, nc = 15, nJ = 0, nH = 0, nHv = 3
pi [*] :=
    african    0.104805
    east_asian 0.264861
    european   0.17273
native_american 0.284324
    south_asian 0.173279
;
```

# Example



```
ampl: reset; model mixtures.mod; data mixtures.dat; option solver conopt, conopt_options 'lsesqp=0'; solve; display pi;
CONOPT 3.17A: lsesqp=0
CONOPT 3.17A: Locally optimal; objective 19.64638424
11 iterations; evals: nf = 18, ng = 11, nc = 18, nJ = 0, nH = 0, nHv = 0
pi [*] :=
    african    0.104805
    east_asian 0.264861
    european   0.17273
native_american 0.284324
    south_asian 0.173279
;

ampl:
ampl: reset; model mixtures.mod; data mixtures.dat; option solver conopt, conopt_options 'lsesqp=1'; solve; display pi;
CONOPT 3.17A: lsesqp=1
CONOPT 3.17A: Locally optimal; objective 19.64638424
10 iterations; evals: nf = 15, ng = 9, nc = 15, nJ = 0, nH = 0, nHv = 3
pi [*] :=
    african    0.104805
    east_asian 0.264861
    european   0.17273
native_american 0.284324
    south_asian 0.173279
;
```

# Example



```
ampl: reset; model mixtures.mod; data mixtures.dat; option solver conopt, conopt_options 'lsesqp=0'; solve; display pi;
CONOPT 3.17A: lsesqp=0
CONOPT 3.17A: Locally optimal; objective 19.64638424
11 iterations; evals: nf = 18, ng = 11, nc = 18, nJ = 0, nH = 0, nHv = 0
pi [*] :=
    african    0.104805
    east_asian 0.264861
    european   0.17273
native_american 0.284324
    south_asian 0.173279
;

ampl:
ampl: reset; model mixtures.mod; data mixtures.dat; option solver conopt, conopt_options 'lsesqp=1'; solve; display pi;
CONOPT 3.17A: lsesqp=1
CONOPT 3.17A: Locally optimal; objective 19.64638424
10 iterations; evals: nf = 15, ng = 9, nc = 15, nJ = 0, nH = 0, nHv = 3
pi [*] :=
    african    0.104805
    east_asian 0.264861
    european   0.17273
native_american 0.284324
    south_asian 0.173279
;
```



# Example



```
ampl: reset; model mixtures.mod; data mixtures.dat; option solver conopt, conopt_options 'lsesqp=0'; solve; display pi;
CONOPT 3.17A: lsesqp=0
CONOPT 3.17A: Locally optimal; objective 19.64638424
11 iterations; evals: nf = 18, ng = 11, nc = 18, nJ = 0, nH = 0, nHv = 0
pi [*] :=
    african    0.104805
    east_asian 0.264861
    european   0.17273
native_american 0.284324
    south_asian 0.173279
;

ampl:
ampl: reset; model mixtures.mod; data mixtures.dat; option solver conopt, conopt_options 'lsesqp=1'; solve; display pi;
CONOPT 3.17A: lsesqp=1
CONOPT 3.17A: Locally optimal; objective 19.64638424
10 iterations; evals: nf = 15, ng = 9, nc = 15, nJ = 0, nH = 0, nHv = 3
pi [*] :=
    african    0.104805
    east_asian 0.264861
    european   0.17273
native_american 0.284324
    south_asian 0.173279
;
```

# Example



```
ampl: reset; model mixtures.mod; data mixtures.dat; option solver conopt, conopt_options 'lsesqp=0'; solve; display pi;
CONOPT 3.17A: lsesqp=0
CONOPT 3.17A: Locally optimal; objective 19.64638424
11 iterations; evals: nf = 18, ng = 11, nc = 18, nJ = 0, nH = 0, nHv = 0
pi [*] :=
    african    0.104805
    east_asian 0.264861
    european   0.17273
native_american 0.284324
    south_asian 0.173279
;

ampl:
ampl: reset; model mixtures.mod; data mixtures.dat; option solver conopt, conopt_options 'lsesqp=1'; solve; display pi;
CONOPT 3.17A: lsesqp=1
CONOPT 3.17A: Locally optimal; objective 19.64638424
10 iterations; evals: nf = 15, ng = 9, nc = 15, nJ = 0, nH = 0, nHv = 3
pi [*] :=
    african    0.104805
    east_asian 0.264861
    european   0.17273
native_american 0.284324
    south_asian 0.173279
;
```

# Gradient Descent Method in Solving Optimization Problems

# Example



```
param N >= 0;    # 20 genetic allele frequency observations
set ancestry;    # 5 global ancestries in the mixtures model

param a_tilde {1..N};    # 20x1 allele frequency from an observed heterogeneous ancestry
param A{1..N, ancestry};    # 20x5 matrix of allele frequencies

var pi {ancestry} >= 0; # propotion value for each ancestry in the mixtures model

minimize HA:
    sum {n in 1..N} (sum {k in ancestry} ((A[n,k] * pi[k]) - a_tilde[n])**2) ;
    # sum of least squares mixtures model

subject to sumtoone:
    sum {k in ancestry} pi[k] = 1; # proportions sum to 1
```

# Example



```

set ancestry := african east_asian european native_american south_asian;
param N := 20;

param a_tilde :=
  1  0.66720042 2 0.17043184    3 0.06935845    4 0.42307358    5 0.77343482
  6  0.20207072 7 0.76939016    8 0.06586837    9 0.32977999   10 0.34011927
 11  0.21841629 12 0.09462661   13 0.11928382   14 0.70109397   15 0.60080761
 16  0.98087062 17 0.88775104   18 0.39677821   19 0.22700411   20 0.62739267;

param A :   african    east_asian    european    native_american    south_asian    :=
  1  0.6352436  0.1385807  0.1665156  0.5659577  0.9041398
  2  0.2625085  0.2629710  0.4181271  0.7178834  0.7896494
  3  0.11973698 0.06742608  0.68754815  0.64800228  0.71269570
  4  0.1713013  0.7879743  0.4803838  0.3807989  0.9183643
  5  0.03612569 0.04574899  0.37353090  0.31219216  0.07174996
  6  0.37142610 0.08522377  0.08455566  0.52003942  0.53396940
  7  0.01295619 0.52020509  0.08344570  0.80204798  0.40205015
  8  0.2384410  0.8258287  0.4895211  0.6896276  0.7545883
  9  0.2051679  0.9381787  0.3756678  0.5106258  0.4657259
 10  0.02585700 0.21247118  0.18287401  0.55115948  0.09243381
 11  0.7119114  0.7760170  0.7929497  0.3597336  0.8863532
 12  0.005123027 0.104136581 0.666181950 0.022666776 0.745053036
 13  0.5524233  0.5716317  0.9709562  0.1433206  0.7371474
 14  0.2896559  0.8975506  0.4844245  0.9709875  0.2204392
 15  0.12560236 0.91613278  0.88561664  0.02557818  0.94784594
 16  0.8064603  0.5860706  0.3659275  0.8404793  0.3409812
 17  0.8327337  0.6031340  0.5502941  0.4890311  0.5284046
 18  0.3112420  0.7053747  0.5472243  0.1844714  0.8127638
 19  0.7856046  0.6434881  0.7781458  0.9617573  0.4538499
 20  0.2991599  0.8962980  0.8208560  0.6753767  0.1596549;

```